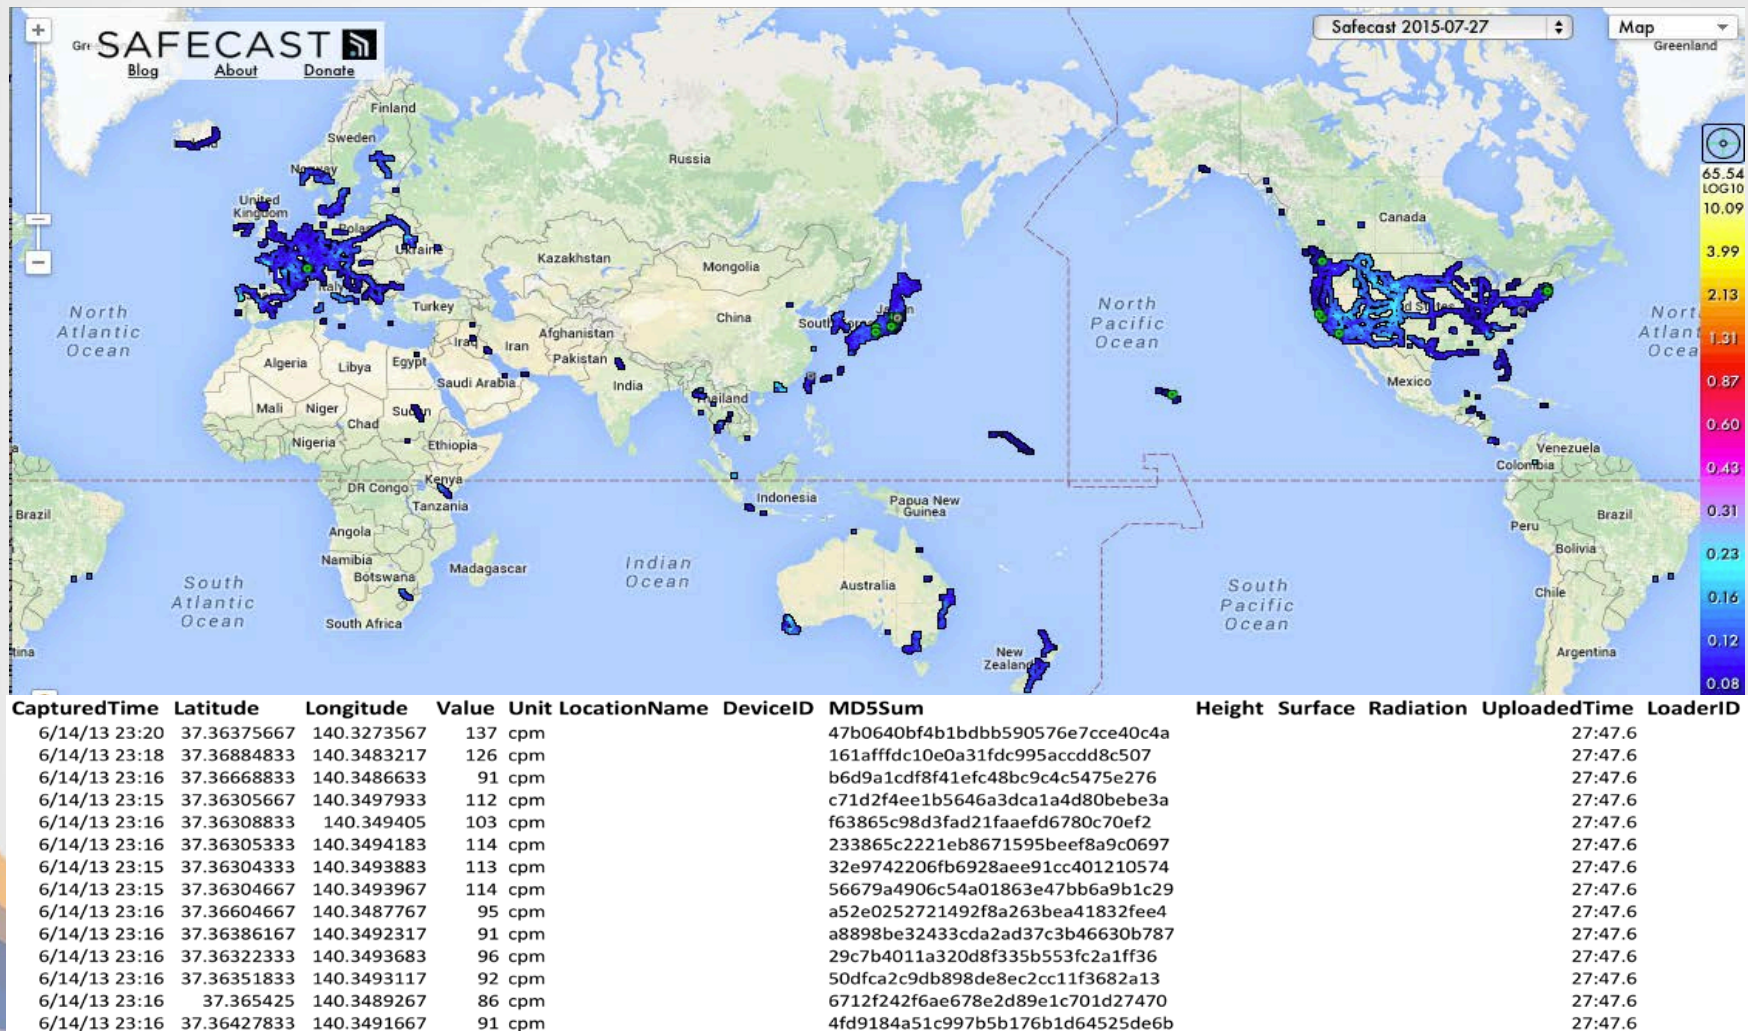


Advanced Geospatial Data Analytics: Pig and Python on Hadoop

Exploring radiation data into HDFS

- Before you begin:
 - The data for the lab is stored in `/gpfs_scratch/geog479/lecture7`
 - Copy the data to your home directory:
 - `cp -r /gpfs_scratch/geog479/lecture7 ~/`

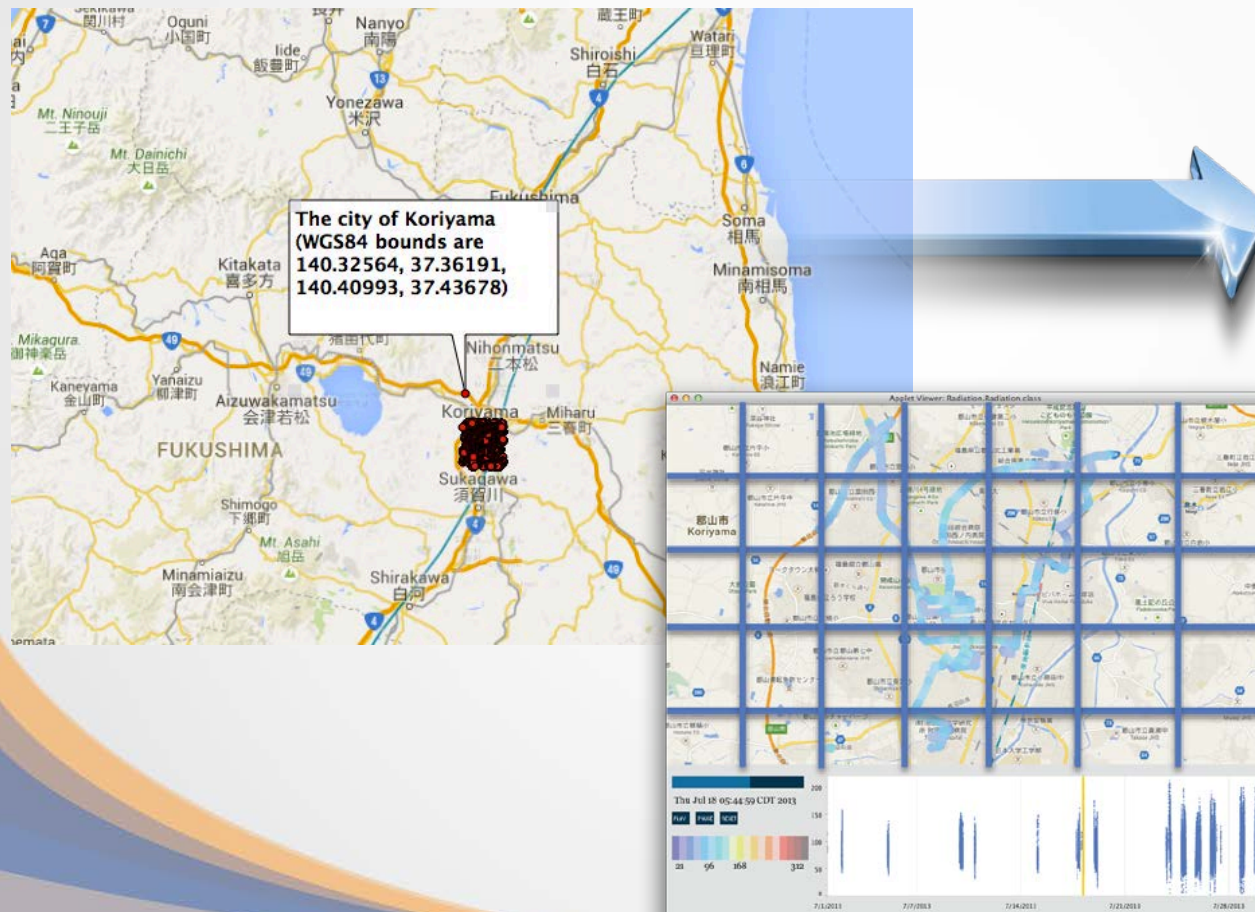
Data 1: Safecast dataset



Case 1: Making a heatmap input data based on radiation dataset

- Objective:
 - Use Pig to navigate through HDFS and explore a dataset.
 - Make a heatmap data for generating a heatmap image file.
- Before you begin:
 - Radiation data is in ~/lecture7/data
 - Load the radiation data to HDFS
 - Optional: check the replication factor we had talked about
 - `hdfs dfs -ls`
 - `hdfs dfs -setrep -w 2 radiation.csv`
 - Now check the HDFS to see what happens

The city of Koriyama, Japan



((0,0), 120.97 cpm)
 ((0,1), 93.77 cpm)
 ((0,2), 108.32 cpm)
 ((0,3), 101.45 cpm)
 ((0,4), 112.87 cpm)
 ((0,5), 98.45 cpm)
 ((0,6), 119.32 cpm)
 ((0,7), 121.21 cpm)

...

Step 1: Making a heatmap data

- Start the Grunt Shell.
 - `$ pig`
- Load the radiation.csv data in HDFS.
 - `$ raw = LOAD 'radiation.csv' USING PigStorage(',') AS (CapturedTime : chararray, Latitude:double, Longitude:double, Value:int, Unit:chararray);`
- **Step 1: Filter out CaptureTime.**
 - `$ step1 = FILTER raw BY CapturedTime is not null AND CapturedTime != 'null' AND CapturedTime != '';`
- **Step 2: select the rows of which unit is cpm.**
 - `$ step2 = FILTER step1 BY $4 MATCHES 'cpm';`
- **Step 3: Filter out the data based on the bounding box of Koriyama, Japan.**
 - `$ step3 = FILTER step2 BY $2 > 140.32564 AND $2 < 140.40993 AND $1 > 37.36191 AND $1 < 37.43678;`

Step 1: Making a heatmap data

- **Step 4: Select the rows of which CapturedTime is August.**
 - `$ step4 = FILTER step3 BY ($0 > '2013-08-01 00:00:00' AND $0 < '2013-09-01 00:00:00');`
- **Step 5: Generate new records to send down the pipeline to the next operator.**
 - `$ step5 = FOREACH step4 GENERATE $0,$1,$2,$3,$4;`
- **Step 6: Generate new records based on latitude and longitude key.**
 - `$ step6 = FOREACH step5 GENERATE $3,(int)((($2 - 140.32564)/0.005) AS logkey,(int)((($1 - 37.36191)/0.005) AS latkey;`
- **Step 7: Collect all records with the same latitude and longitude key.**
 - `$ step6_group = GROUP step6 BY (latkey,logkey);`

Step 1: Making a heatmap data

- **Step 8: Calculate the average of radiation levels per each cell.**
 - `$ step6_avg = FOREACH step6_group GENERATE group, AVG(step6.Value);`
- **Step 9: Sort data based on the group.**
 - `$ result = ORDER step6_avg BY group;`
- **Step 10: Send the output to a folder in HDFS.**
 - `$ Store result into 'heatmap_output';`
- **Step 11: Log out the Grunt shell**
 - `$ quit (or by Ctr + c)`

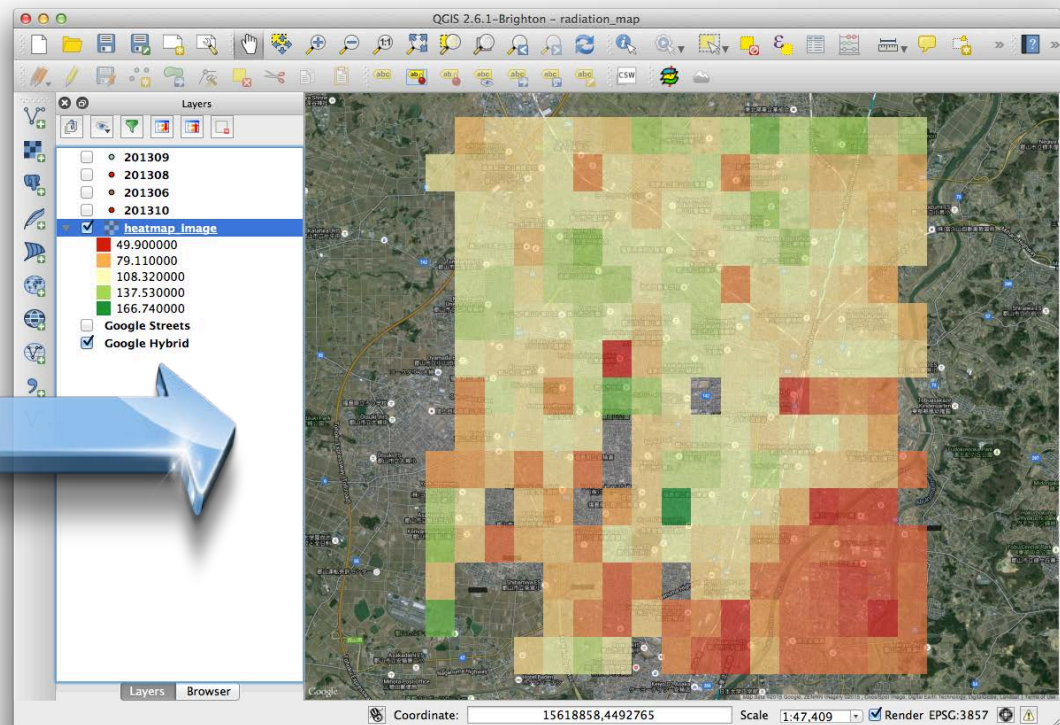
Step 2: Making a heatmap

- Step 13: Take a source directory as input and concatenate files into the destination local file
 - `$ hadoop fs -getmerge heatmap_output ~/lecture7/demo1/heatmap.txt`
- Check the content of heatmap.txt
 - `$ tail heatmap.txt`
 - `python demo1.py`

```
[cgtrn20@cg-hm06 Exe2]$ tail heatmap.txt
(14,7) 130.57432432432432
(14,8) 120.9330985915493
(14,9) 101.0
(14,10) 119.0
(14,11) 145.14545454545456
(14,12) 123.13559322033899
(14,13) 134.55882352941177
(14,14) 137.2
(14,15) 103.01149425287356
(14,16) 122.54545454545455
[cgtrn20@cg-hm06 Exe2]$
```

Visualize the generated heatmap

```
[cgtrn20@cg-hm06 Exe2]$ tail heatmap.txt  
(14,7) 130.57432432432432  
(14,8) 120.9330985915493  
(14,9) 101.0  
(14,10) 119.0  
(14,11) 145.14545454545456  
(14,12) 123.13559322033899  
(14,13) 134.55882352941177  
(14,14) 137.2  
(14,15) 103.01149425287356  
(14,16) 122.54545454545455  
[cgtrn20@cg-hm06 Exe2]$
```





Analyzing data with Pig UDF using Python

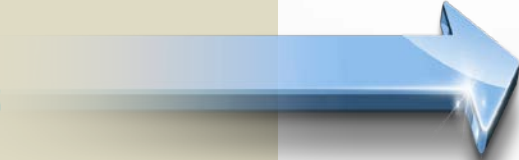
Python in ROGER

- There are multiple versions of Python installed in ROGER
 - List all the available Python environment in ROGER
- In particular, the Python environment comes with Anaconda is used in this course
 - scipy, numpy, Pandas, Sciki-learn, gdal-python, etc.
 - IPython
 - IPython notebook

Case 2: Generating a heatmap after removing outliers using Pig UDF

- Objective:
 - Make a heatmap image file (Ascii format) after removing outliers per each cell.
 - Compare the result with the one from the previous exercise.

```
[cgtrn20@cg-hm06 Exe2]$ head heatmap.txt
(0,3) 96.57894736842105
(0,4) 107.69850402761796
(0,5) 125.0909090909091
(0,6) 105.25
(0,8) 71.88888888888889
(0,9) 61.470588235294116
(0,10) 54.3
(0,11) 87.21832884097034
(0,12) 66.65137614678899
(0,13) 67.17173252279635
[cgtrn20@cg-hm06 Exe2]$
```



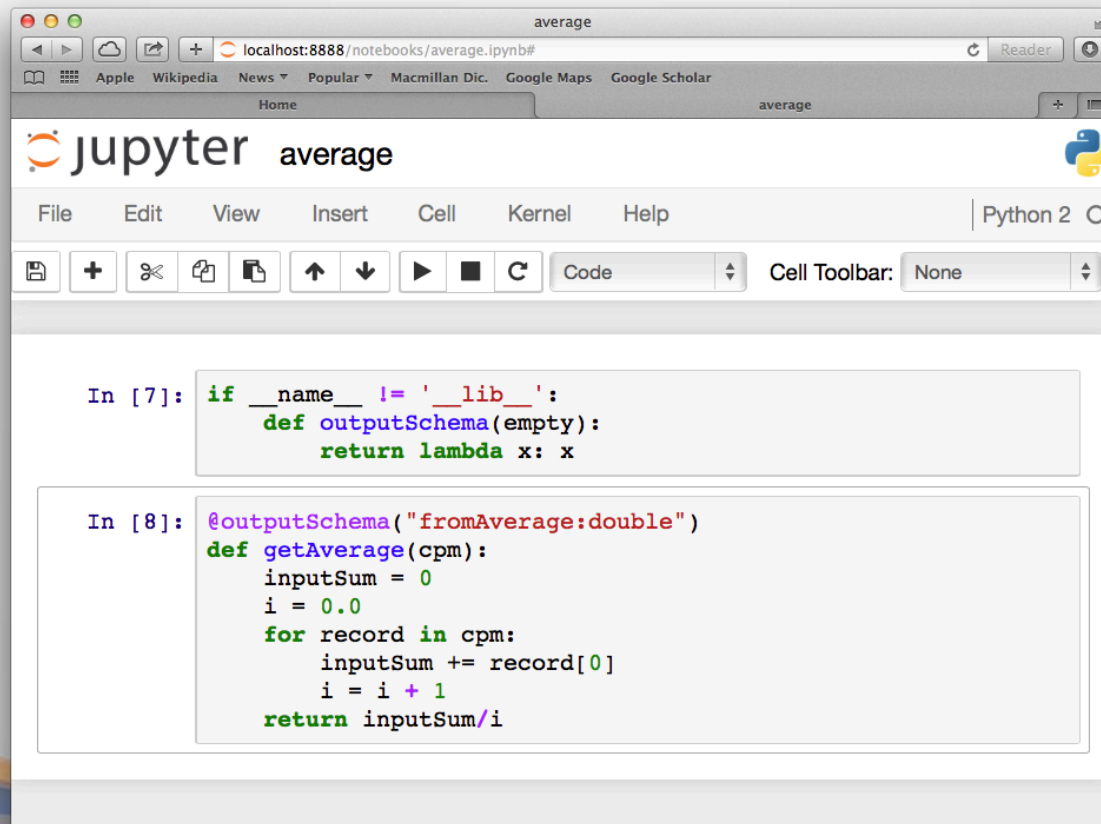
```
[cgtrn20@cg-hm06 Demo2]$ head heatmap2.txt
(0,3) 96.57894736842105
(0,4) 98.75354838709677
(0,5) 125.0909090909091
(0,6) 105.25
(0,8) 71.88888888888889
(0,9) 61.470588235294116
(0,10) 54.3
(0,11) 86.8070652173913
(0,12) 66.65137614678899
(0,13) 65.84101748807632
[cgtrn20@cg-hm06 Demo2]$
```


Pig User Defined Functions (UDFs)

- Alternatives to Hadoop streaming for running Python on a cluster
 - Pig User Defined Functions (UDFs): Invoking Python code from within a Pig script
 - Pig Streaming: Using the Pig STREAM command
- Combine Pig's operators with their own or others' code via UDFs
- Pig provides extensive support for user defined functions (UDFs) as a way to specify custom processing
 - Java
 - Jython: a Python implementation that runs on a JVM
 - Python (Python UDFs do not work on Hadoop 2.x)
 - Javascript, Ruby and Groovy (experimental)

Writing a Pig UDF in Python

- A UDF written in Python can be any Python function.
- Annotate the function using `@outputSchema` so Pig can determine what the data type of the return value is.



```
In [7]: if __name__ != '__lib__':
        def outputSchema(empty):
            return lambda x: x

In [8]: @outputSchema("fromAverage:double")
        def getAverage(cpm):
            inputSum = 0
            i = 0.0
            for record in cpm:
                inputSum += record[0]
                i = i + 1
            return inputSum/i
```

Invoking a UDF

- Before invoking a Python UDF, we need to register the text file that contains the function definition.
 - register 'udfs.py' using `jython` as `udfs`;
 - `A = LOAD 'work' AS (department:chararray,name:chararray,salary:double);`
 - `B = GROUP A BY department;`
 - `C = FOREACH B GENERATE group, udfs.getAverage(A.Value) ;`
- By incorporating the UDF in the Pig script, we can take advantage of Pig's ability to manipulate and group our big data, while also taking advantage of Python's powerful but simple processing and logic.

Invoking a UDF

- Work on the directory where you have the scripts ready
 - `cd ~/lecture7/demo2`
- Run the pig script and see the results
 - What are the precedures?

Case 3: Implementing K-Means clustering on Hadoop for 311 Data

- Objective:
 - Run a Pig script that streams data through a Python script.
 - Use Scikit-Learn to generate a K-Means clustering for 311 Data.
- Before you begin:
 - You need to load module anaconda, which contains the k-mean clustering library you need
 - module load ?
 - Work on the directory where you have the scripts
 - `cd ~/lecture7/demo3`

Pig STREAM Command

- The Pig STREAM command sends data to an external script or program. The syntax looks like:
 - `alias = STREAM alias [, alias ...] THROUGH {`command` | cmd_alias}`
`[AS schema] ;`
 - E.g., `x = LOAD '311_data';`
`result = STREAM x THROUGH `kmeans`;`
- Defining a streaming alias
 - Use the DEFINE command to assign an alias to a streaming command.
 - `DEFINE kmeans `kmeans311.py` SHIP ('kmeans311.py')`
 - Specify files that are needed on the worker nodes when using Pig streaming.

311 Dataset

- 311 is for non-emergency services.
- Improved customer service for city departments and governmental agencies
- Service requests:
 - Graffiti Removal, Vacant and Abandoned Buildings, Pot Holes, etc.
- ID, Date, Request service, Latitude, Longitude

13-01771357	11/26/13	Sanitation Code Violation	41.7641222	-87.655594
13-01652015	11/4/13	Sanitation Code Violation	41.7549537	-87.565082
13-01759727	11/23/13	Sanitation Code Violation	41.7553423	-87.566302
13-01733479	11/19/13	Sanitation Code Violation	41.7532822	-87.562626
13-01782724	11/28/13	Sanitation Code Violation	41.7506552	-87.600318
13-01639956	11/1/13	Pot Hole in Street	41.9388159	-87.756114
13-01639988	11/1/13	Pot Hole in Street	41.8779797	-87.639706
13-01640011	11/1/13	Pot Hole in Street	41.887046	-87.613691
13-01640019	11/1/13	Pot Hole in Street	41.9376795	-87.688288
13-01640044	11/1/13	Pot Hole in Street	41.8767974	-87.633256

Python code to implement the K-Means clustering algorithm

Define the run_kmeans function

```
def run_kmeans(dataset, max_iterations=100, num_clusters=10, num_seeds=10):
    geo_locations = DataFrame(dataset, columns=['lat', 'lon'])
    #Create a KMeans object
    km = KMeans(n_clusters=num_clusters, init='k-means++',
               max_iter=max_iterations, n_init=num_seeds, verbose=0)
    #Compute cluster centers and predict cluster index
    clusters = km.fit_predict(geo_locations)
    #Return the results in a new DataFrame
    result = DataFrame(dataset, columns = ["id", "time", "request", "lat", "lon"])
    result["cluster_id"] = clusters
    return result
```

Process the input data (line by line)

```
dataset = []
for line in fileinput.input():
    line = line.strip()
    fields = line.split('\t')
    dataset.append({"id": fields[0], "time": fields[1], "request": fields[2],
                   "lat": fields[3], "lon": fields[4]})
```

Pass the data to the clustering function

```
clusters = run_kmeans(DataFrame(dataset))
print(DataFrame.to_string(clusters, header = False, index = False))
```

Stream the data through a clustering algorithm using Pig streaming

- `define kmeans311 `/sw/anaconda/bin/python kmeans311.py` ship ('kmeans311.py');`
- `raw = LOAD 'cybergis/311_data.csv' USING PigStorage(',') AS (id:chararray,time:chararray,request:chararray,Lat:double,Lon:double);`
- `step0 = FILTER raw BY time != 'null' AND time != '';`
- `step1 = FILTER step0 BY $0 is not null;`
- `step2 = FILTER step1 BY $1 is not null;`
- `step3 = FILTER step2 BY $2 is not null and $2 MATCHES 'Alley Light Out';`
- `step4 = FOREACH step3 GENERATE $0,$1,$2,$3,$4;`
- `result = stream step4 through kmeans311;`
- `store result into 'kmeans_output';`

K-Means Clustering

- Step 1: Put the data in HDFS
 - `$hdfs dfs -copyFromLocal 311_data.csv`
- Step 2: Run the Pig Script
 - `$pig kmeans311.pig`
 - Save the result to the local directory
 - `$hdfs dfs -getmerge kmeans_output kmeans.csv`
 - Visualize the results as points in ArcGIS or QGIS
 - Take advantage of MobaXterm

K-Means Clustering

- Step 3: View the result
 - \$tail kmeans.csv
 - K-Means clustering for Graffiti removal requests from 01/15/2015 to 08/01/2014.

