

# Distributed Spatial Database

Dr. Junjun Yin

*CyberGIS Center for Advanced Digital and Spatial Studies  
Department of Geography and Geographic Information Science  
National Center for Supercomputing Applications  
University of Illinois at Urbana-Champaign, IL, 61801, USA  
[jyn@illinois.edu](mailto:jyn@illinois.edu)*

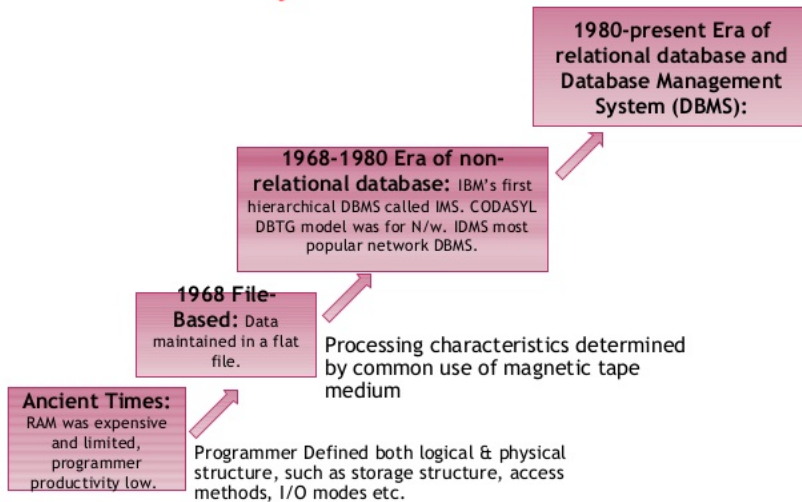
*February 9, 2016*

# Outline

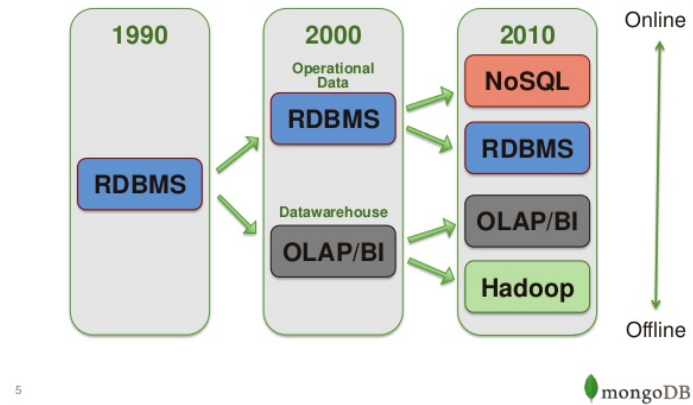
- What is “distributed databases”
  - File system vs. database
- Relational database vs NoSQL database
  - Relational Database Management System (RDBMS)
  - Spatial Database Management System (SDBMS)
  - Non-relational database (NoSQL)
- Big Data requirements
  - Variety and data structure
  - Velocity and concurrency
- Data management, storage and query performance

# Evolution of Databases

## Database Systems: A Brief Timeline





## The Evolution of Databases



## The early days














- Navigational Databases
  - Objects are found by following references from other objects
- Network and Hierarchical Databases
  - Data is organized into a network or a tree-like structure



Source:  
<http://www.slideshare.net/wlaforest/why-nosql-and-mongodb-for-big-data>

# Evolution of Databases



Document Database	Graph Databases
 Couchbase  MarkLogic  mongoDB	 Neo4j  InfiniteGraph <small>The Distributed Graph Database</small>
Wide Column Stores	Key-Value Databases
 redis  amazon DynamoDB  AEROSPIKE  riak	 HYPERTABLE  accumulo  Cassandra  APACHE HBASE Amazon SimpleDB

@cloudtxt <http://www.aryannava.com>

Source: <http://aryannava.com/category/nosql/>

# Document-oriented Database

```
{  
  "_id" : ObjectId("5660d0d92a5d502680eb616f"),  
  "loc" : {  
    "type" : "Point",  
    "coordinates" : [  
      -87.9609112,  
      42.1526175  
    ]  
  },  
  "message" : "3 more!!!\n",  
  "user_id" : "232659889",  
  "time" : "Fri Jul 04 20:55:28 CDT 2014"  
}
```

- JSON (JavaScript Object Notation) based data format

# Column-oriented Database

Row Oriented Database

<u>date</u>	<u>price</u>	<u>size</u>
2011-01-20	10.1	10
2011-01-21	10.3	20
2011-01-22	10.5	40
2011-01-23	10.4	5
2011-01-24	11.2	55
2011-01-25	11.4	66
...	...	...
2013-03-31	17.3	100

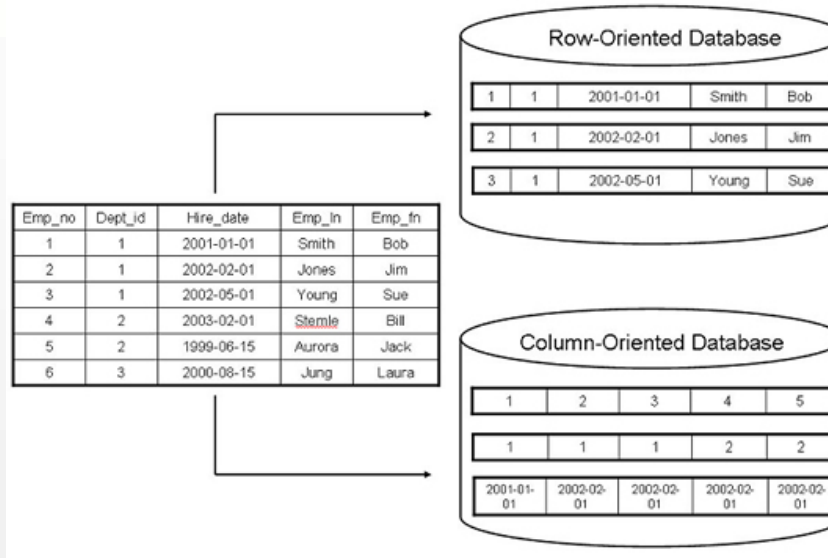
Table of Data

<u>date</u>	<u>price</u>	<u>size</u>
2011-01-20	10.1	10
2011-01-21	10.3	20
2011-01-22	10.5	40
2011-01-23	10.4	5
2011-01-24	11.2	55
2011-01-25	11.4	66
...	...	...
2013-03-31	17.3	100

Column Oriented Database

<u>date</u>	<u>price</u>	<u>size</u>
2011-01-20	10.1	10
2011-01-21	10.3	20
2011-01-22	10.5	40
2011-01-23	10.4	5
2011-01-24	11.2	55
2011-01-25	11.4	66
...	...	...
2013-03-31	17.3	100

Source:  
<http://www.timestored.com/time-series-data/what-is-a-column-oriented-database>





# Relational vs. NoSQL


- Relational databases
  - Well designed data schemes
  - Well organized data structures
  - Row-based
- NoSQL databases
  - Document-oriented
  - Column-oriented
  - <key, value> pair-based
  - Graph-based

# Spatial Database

- Geometry
  - Point, polyline, polygon, and 3D geospatial objects
- Attributes
- Spatial indexing
  - R-tree ( $R^*$ -tree), 3D R-tree
  - Quad-tree
- Topological operations
  - Overlapping
  - Intersection
  - K-nearest neighbors
  - Within
  - ...



# Distributed Databases



# Why distributed database

- Your single machine runs out of storage!
  - Easy scalability
- Avoids single point of failure by distributing data among multiple machines.
  - Will not lose any data in case a machine fails.
  - Will not lose the ability to query the data in case a machine fails.
  - Machines that are performance bottlenecks will not affect the overall performance of the system.

# Why distributed database

- Reduce query time by “intelligently” distribute data among multiple nodes.
  - A query should look into a smaller subset of data.
  - In write operations the index should be updated for a smaller set of data.
- Distribute workload on multiple nodes in case of simultaneous access.
  - Re-route requests to multiple machines (courtesy of replication).
- Incorporate organizational structure.
  - Each group wants to host their own data!

# Challenges

- Require new skills and knowledge for many database admins.
- Relatively new era. Still lots of development is being made.
- Issues with integrity and consistency of the database.
- More issues on ensuring the security of the database cluster.
- Require new database design techniques.

# Concepts

- Replication: System maintains multiple identical copies of the data.
  - Increase latency on update.
- Partitioning: The dataset is divided into multiple smaller datasets based on a pre-defined function.
  - Balancing the load on different chunks of data should be considered.
- Consistency: any transaction must only affect data in allowed ways.

# Consistency

**X : John's balance**



X: 1000\$



X: 950\$



X: 1020\$



X: 1008\$



# Types of Consistency

- **Strong consistency:** Every client can observe one consistent stage.
- **Weak consistency:** relaxing some constraints on the strong consistency to gain better performance and availability.

# CAP Theorem

It is impossible for a distributed computer system to simultaneously provide all of these 3 guarantees:

- Consistency
- Availability
- Partition tolerance
  - The system can continue its operation despite partitioning due to failures.


# Examples of NoSQL Distributed Databases



mongoDB



redis



*cassandra*

# MongoDB

- Document-oriented database.
  - Store records in a json-like format.
- Easy-to-use.
- Supports replication and partitioning (sharding).
- Support primary and secondary indexes.
- Support aggregation queries through:
  - Javascript-based MapReduce structure which can support arbitrary large data but it is slow.
  - C++ based aggregation framework which is limited in size but very fast.

# Apache Cassandra

- Highly decentralized architecture.
- Read/write throughput both scale linearly as the new machines are added.
- Very fast write performance.
- The admin can tune the consistency level based on the application.
- Most popular option in the industry for massive and highly distributed datasets.
- Easy integration with Hadoop, Hive, etc.
- Limited support for geospatial operations.

# Data Model

Partitioning Key

Clustering Key

Row ID	Field 1   Field 2	Field 1   Field 2	Field 1   Field 2
	Other fields	Other fields	Other fields

Bus Station	Time   Bus Id	Time   Bus Id	Time   Bus Id
	# of passenger	# of passenger	# of passenger



# Query Language: CQL

```
INSERT INTO BUSDATA (station, time, id, passengers) VALUES (100,  
201504041000, 13222, 22);
```

```
SELECT * FROM BUSDATA WHERE station = 100 AND time >=  
201510010000;
```

```
UPDATE BUSDATA SET passengers = passengers + 5 WHERE station =  
100 AND time <= 201511010000 AND time >=201510010000;
```

# Redis

- In memory key-value stores.
- Supports limited operations but it is significantly faster for smaller datasets.
- Supports replications.
- Provide multiple methods to ensure persistency of the data.
- Main supported data types:
  - Lists
  - Sets and Sorted Sets.
  - Hash Tables.

# Basic Commands

SET "13132" "John Smith"

GET "13132"

"John Smith"

ZADD "arrivals" 201410100105 "american airlines 3322"

ZRANGEBYSCORE "arrivals" 201410010000 201410052359

HSET "employees" "13132" "John Smith"

HGET "employees" "13132"

# Geospatial Commands

```
GEOADD Illinois -87.6847 41.8369 "Chicago"
```

```
GEOADD Illinois -88.2728 40.1150 "Champaign"
```

```
GEODIST Illinois Champaign Chicago mi
```

```
GEORADIUS Illinois -88.2000 40.3442 200 km WITHDIST
```

Details: <http://redis.io/commands/georadius>

# Performance Comparison

## **Solving Big Data Challenges for Enterprise Application Performance Management**

**Mainly from University of Toronto  
VLDB 2012**

- Benchmark the databases with different workloads (read-heavy, write-heavy, mixed).
- Benchmark the database with different metrics (latency (sec) and throughput (ops/sec)).

# Performance Comparison

- Cassandra is a good choice for applications with high insertion rate and applications which requires great scalability.
- Redis performs significantly faster than other database systems in read-heavy loads, however it does not scale well. In addition it comes with a limitation on the size of data.
- MongoDB works fine for smaller clusters, but shows significant performance drop as we move to larger clusters.



# Case Study: Spatial Query Engine


- Establish a generic and scalable query framework for geospatial big data on cyberinfrastructure.
- Support query patterns in a scalable way.
- Use Redis to store the indexing structure in memory and the actual data on Cassandra database.
- Design for integrated spatiotemporal queries.
- Arbitrary input and query shapes.

# Case Study: Spatial Query Engine

## Polygon Indexing



## Point Indexing



# Experiments

- 2 weeks of geo-tagged tweets in US
- Rectangular query shape
- 100 queries per level (point and 3 polygon levels)
- Temporal length: one week
- Maximum width/height in miles
  - Point: 2.0
  - Zipcodes: 16.0, County: 64.0, State: 256

# Result I

*Database write and read -- local performance (ms)*

Operation/Node	Average for Point data	Average for Polygon Data
Local Write Latency	0.0445	0.984
Local Read Latency	21.738	3.122

## Result II

*End-to-end query performance(ms)*

Query Type/Latency	Mean	Average	90% Percentile
Polygon Level 0	82	85	160
Polygon Level 1	78	80	142
Polygon Level 2	158	190	322
Point	75	121	200

# MongoDB in CyberGIS OpenStack

- Managing MongoDB in OpenStack
  - Create a new instance(s) in OpenStack to host MongoDB server (clusters)
  - For this class, we have created a MongoDB instance
  - Use a GUI client to access MongoDB
- Using MongoDB
  - Choose your favorite programming language to access the databases
  - Using command line interface:  
mongo --host **hostserver\_name:port (default: 27017)/database\_name**  
mongo --host 141.142.168.54  
show dbs  
use **<name>** to create a new database



# Geo-located Tweets query in Chicago

- Getting started with MongoDB in OpenStack
  - Choose your programming language and setup appropriate drivers
  - List of commands and guidelines:
    - <https://docs.mongodb.org/manual/core/crud-introduction/>
- Prepare geo-located Tweets as documents to MongoDB
  - MongoDB uses JSON objects to store the data
  - Convert your data structure to JSON
  - Covert geospatial objects to GeoJSON
  - They are essentially <key, value> pairs but with hierarchical structure

# Geo-located Tweets query in Chicago

- A GeoJSON object

```
{  
  "type": "Feature",  
  "geometry": {  
    "type": "Point",  
    "coordinates": [125.6, 10.1]  
  },  
  "properties": {  
    "name": "Dinagat Islands"  
  }  
}
```

- Geo-located tweets

- Using Pig script to extract <User\_id, latitude, longitude, timestamp, content>

# Import Twitter Data into MongoDB

- `from pymongo import MongoClient`
- `import datetime`
- `import json`
- `def main():`
- `client = MongoClient('mongodb://141.142.168.54:27017')`
- `db = client.ddbs`
- `collection = db['chicago']`
- `mFile = open("2014_12_25_chi.txt", "rb")`
- `for mLine in mFile:`
- `# print mLine`
- `[uid, lat, lng, tm, msg] = mLine.split("\t")`
- `geoString = {'user_id': uid, 'time': tm, 'message': msg, 'loc': {'type': 'Point',`
- `'coordinates': [lng, lat]}}`
- `print geoString`
- `db.chicago.insert(geoString)`
- `if __name__ == '__main__':`
- `main()`

# Geo-located Tweets query in Chicago

- Create spatial index to speed up database querying
  - `db.chicago.createIndex( { loc : "2dsphere" } )`
- Find the geo-located tweets within the search distances from a specified location  
`db.chicago.find({loc: {$nearSphere: {$geometry: {type : "Point",coordinates : [-87.639160, 41.878628]}, $minDistance: 1000, $maxDistance: 5000}}})`
- Find the geo-located tweets within a specified region (polygon)  
`db.chicago.find({loc:{$geoWithin: { $geometry: {type : "Polygon",  
coordinates: [ [ [-87.639546, 41.878053], [87.639643, 41.879506], [-87.638602,41.879363], [-87.638270,41.878045], [-87.639546, 41.878053]]]]}}})`

# Geo-located Tweets query in Chicago

- Sort the query result based on distance

```
db.chicago.find({loc: { $nearSphere: { $geometry: { type : "Point",  
coordinates : [-87.639160, 41.878628]},          $minDistance: 1000,  
$maxDistance: 5000}}}).sort({"loc":1})
```



# Pick up coordinates from Google Maps

